

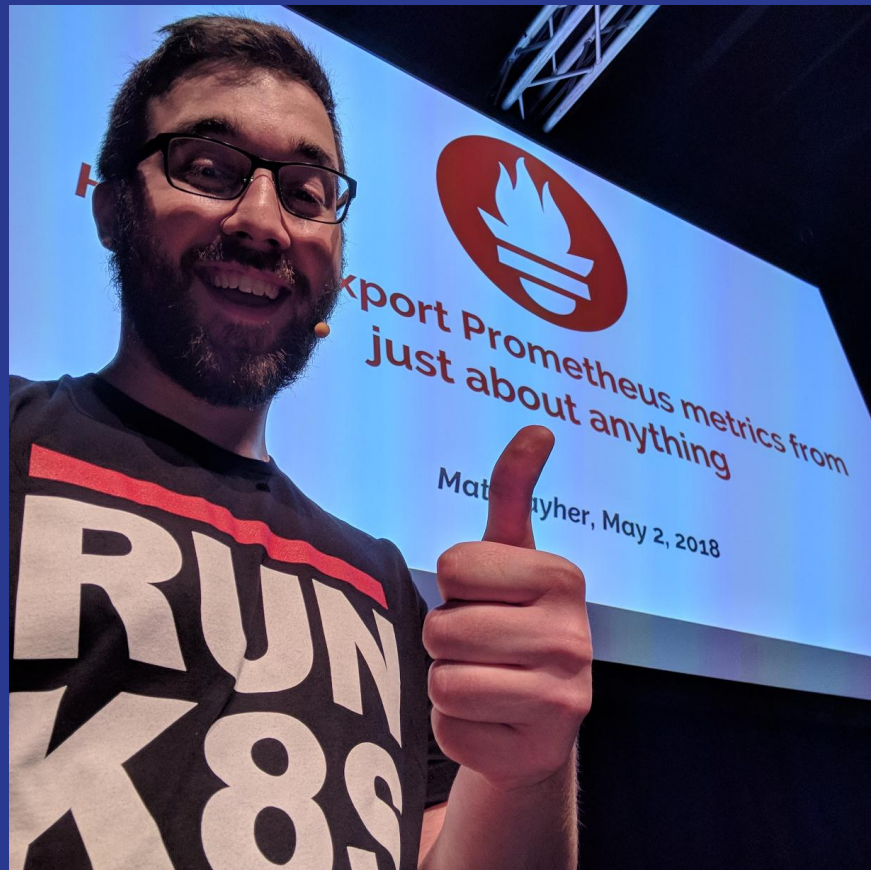


# Implementing a Network Protocol in Go

Matt Layher, August 29, 2018

# Matt Layher

- Engineer at DigitalOcean
- GitHub + Twitter: @mdlayher
- [github.com/mdlayher/talks](https://github.com/mdlayher/talks)



# Intro

- The IPv6 **Neighbor Discovery Protocol (NDP)** is our focus for today
- IPv6 is an important step for the Internet
- Go can be used for many low-level networking applications




# Agenda


- Introduction to IPv6 and NDP
- Using NDP with Go: [github.com/mdlayher/ndp](https://github.com/mdlayher/ndp)
- Building and testing network protocol packages



# Introduction to IPv6




How many  
audience members  
are familiar with  
IPv6?



How many  
audience members  
use IPv6 at home?

Native: 20.43% 6to4/Teredo: 0.00% Total IPv6: 20.44% | Aug 8, 2018



User IPv6 adoption on Google services: <https://www.google.com/intl/en/ipv6/statistics.html>



# What is IPv6?

- The next generation of the Internet Protocol
- Draft standard in [RFC 2460](#) (December 1998!)
- 128 bit IP addresses, huge improvement over 32 bit IPv4 addresses:
  - **IPv6:**  $2^{128}$  addresses, but not all used for hosts
  - IPv4:  $2^{32}$  addresses



# How is IPv6 different from IPv4?

- Wire format simplified, more extensible
- Residential ISPs can offer entire IPv6 prefixes instead of 1 IPv4 address:
  - **IPv6:** 2001:db8:abcd:ffff::/64:  **$2^{64}$  addresses**
  - IPv4: 192.0.2.10/32: 1 address



# IPv6 tips and tricks

- Many shell utilities have a “-6” flag to use IPv6
- A useful website for testing IPv6 configuration: [ipv6-test.com](https://ipv6-test.com)
- My favorite ping target:

```
$ ping6 2600::
```

```
PING 2600::(2600::) 56 data bytes
```

```
64 bytes from 2600::: icmp_seq=1 ttl=48 time=56.9 ms
```



# Introduction to NDP

# What is NDP?

- Effectively the IPv6 equivalent to IPv4's ARP
- Runs on top of IPv6 + ICMPv6 with link-local addresses: fe80::/10
- Used to ask a network neighbor for its MAC address using IPv6 address
  - **A:** Who has "**B**"? Tell "**A**".
  - **B:** "**B**" is at "**04:18:d6:a1:ce:b7**".



# ICMPv6 header

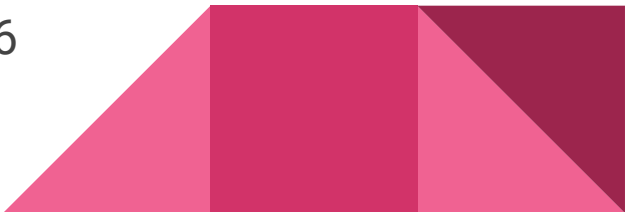
Type (8 bits)	Code (8 bits)	Checksum (16 bits)
Data (N bytes)		

# NDP vs ARP

	<b>NDP</b>	<b>ARP</b>
Transport	IPv6 + ICMPv6	Ethernet
Traffic type	Multicast	Broadcast
Options	Yes	No
Router discovery	Yes	No
Address assignment	Yes	No



# IPv6 and NDP's big advantage

- **DHCP is not usually necessary** to configure globally-routable IPv6 addresses:
    - **Stateless Address Autoconfiguration (SLAAC) via NDP**
      - No DHCPv6 required whatsoever
    - SLAAC + Stateless DHCPv6
      - Addresses via SLAAC, more configuration via DHCPv6
    - Stateful DHCPv6
      - All addresses and information from DHCPv6
- 



# SLAAC via NDP

- SLAAC uses NDP router advertisements to provide IPv6 prefix information
  - “**A**” sends a router solicitation
  - “**R**” sends a router advertisement:
    - Prefix “**P::/64**”, use SLAAC, valid for 24 hours
  - “**P:76d4:35ff:fee7:cbc4**” computed and assigned





# NDP and Go

# NDP and Go


- Your operating system usually handles NDP; why is it useful for Go programs?
- [github.com/mdlayher/ndp](https://github.com/mdlayher/ndp): Go package for using NDP
  - [MetalLB](#): implements IPv6 Layer 2 mode for Kubernetes load balancer
  - [cmd/ndp](#): tool for generating and capturing NDP traffic
  - DigitalOcean internal use: responding to neighbor/router solicitations



# Package ndp overview

- Primary types:
  - **ndp.Message** interface: marshaling/unmarshaling of NDP messages
  - **ndp.Option** interface: marshaling/unmarshaling of NDP options
  - **ndp.Conn** struct: manage ICMPv6 connection, read/write ndp.Messages





How do we go from  
bytes to a complete  
NDP package?



From bytes to messages

# NDP message basics

- ICMPv6 header determines which NDP message is used
  - Type specifies NDP message, Code always 0
- Initial NDP messages and options defined in [RFC 4861](#)
  - Fixed length messages, variable options



# Parsing bytes

- An ICMPv6 header will always precede an NDP message
- NDP messages on their own are not useful without the ICMPv6 header
- Exporting marshal/unmarshal methods bloats the API and GoDoc
  - **Solution:** add functions which always add/remove the ICMPv6 header





# ndp.Message interface

- Exported Type method for documentation, but other methods unexported

```
// A Message is a Neighbor Discovery Protocol message.
```

```
type Message interface {
```

```
    // Type specifies the ICMPv6 type for a Message.
```

```
    Type() ipv6.ICMPType
```

```
    // Called via MarshalMessage and ParseMessage.
```

```
    marshal() ([]byte, error)
```

```
    unmarshal(b []byte) error
```

```
}
```



# ndp.ParseMessage function

- Bounds checking validation, determine concrete type, continue parsing

```
func ParseMessage(b []byte) (Message, error) {  
    // Bounds check!!!  
  
    // Determine ndp.Message from ICMPv6 type.  
  
    // Unmarshal ICMPv6 data into ndp.Message implementation.  
}
```



# Bounds checking

- When using slice elements, you must perform bounds checks to avoid panics


```
// The ICMPv6 header is fixed length.  
const icmpLen = 4  
if len(b) < icmpLen {  
    return nil, io.ErrUnexpectedEOF  
}
```



# Determining `ndp.Message` type

- Use a switch to choose the right interface implementation

```
// Select the correct ndp.Message type based on ICMPv6 header.  
var m Message  
switch t := ipv6.ICMPType(b[0]); t {  
case ipv6.ICMPTypeNeighborSolicitation:  
    m = new(NeighborSolicitation)  
default:  
    return nil, fmt.Errorf("ndp: unrecognized ICMPv6 type: %d", t)  
}
```



# Unmarshal the `ndp.Message` implementation

- Call into the type's methods to do the rest of the work; skipping the header

```
// Unmarshal remaining bytes into correct ndp.Message type.  
if err := m.unmarshal(b[icmpLen:]); err != nil {  
    return nil, err  
}
```



# Parsing messages

- Using `ndp.ParseMessage`, it's easy to parse `ndp.Message` types
- Concise API: one parsing function
- **Correctness and simplicity first**, performance optimizations later



# Our first `ndp.Message`

- Neighbor Solicitation (NS) messages ask a machine for its MAC address
- For now, `ndp.Option` is unimplemented

```
// An Option is a Neighbor Discovery Protocol option.  
type Option interface {  
    // TODO!  
}
```



# ICMPv6 header + NDP NS message

<b>135</b>	<b>0</b>	Checksum (16 bits)
Reserved (32 bits)		
Target IPv6 Address (128 bits)		
Options (N bytes)		



# ndp.NeighborSolicitation type

- Mimic structure defined by RFC, use doc comments to provide references

```
// A NeighborSolicitation is a Neighbor Solicitation message as  
// described in RFC 4861, Section 4.3.
```

```
type NeighborSolicitation struct {  
    TargetAddress net.IP  
    Options       []Option  
}
```



# Checking for IPv4 and IPv6 addresses


- `net.IP` can contain IPv4, IPv6, or totally invalid IP addresses
  - A combination of `To4` and `To16` methods determine the actual type
  - Something I'd love to see improved upon in Go 2
    - `net.IP` interface? `net.IPv4` and `net.IPv6` types?



# checkIPv6 function

```
// checkIPv6 verifies that ip is an IPv6 address.
func checkIPv6(ip net.IP) error {
    // To16 returns nil when ip is not a valid IPv4/IPv6 address.
    //
    // To4 returns non-nil when ip is an IPv4 address.
    if ip.To16() == nil || ip.To4() != nil {
        return fmt.Errorf("ndp: invalid IPv6 address: %q",
            ip.String())
    }

    return nil
}
```



# ndp.NeighborSolicitation unmarshaling

- Validate incoming bytes, replace the structure all at once

```
func (ns *NeighborSolicitation) unmarshal(b []byte) error {  
    // Bounds checking!!!  
  
    // Validation  
  
    // Replacing contents of the NeighborSolicitation  
}
```



# Validating byte inputs

- Ensure that field values make sense, typically using rules defined by RFC

```
// Skip reserved area.  
addr := b[4:nsLen]  
if err := checkIPv6(addr); err != nil {  
    return err  
}
```



# Replacing the structure while unmarshaling

- Dereference pointer and replace contents with completed structure
- Always make a copy of data from the input slice; don't assume it's safe to retain

```
*ns = NeighborSolicitation{  
    TargetAddress: make(net.IP, net.IPv6len),  
    Options:      options,  
}
```

```
copy(ns.TargetAddress, addr)
```





From messages to bytes

# Marshaling messages

- An ICMPv6 header will always precede an NDP message
- NDP messages on their own are not useful without the ICMPv6 header
- Do the parsing operation in reverse





# ndp.MarshalMessage function

- Marshal an ndp.Message into binary, prepend ICMPv6 header

```
func MarshalMessage(m Message) ([]byte, error) {  
    // Call m's marshal method  
  
    // Pack bytes into an ICMPv6 header  
}
```



# Marshaling `ndp.Messages`

- **Simplicity wins**; allocating is okay until your performance needs are not met

```
mb, err := m.marshal()  
if err != nil {  
    return nil, err  
}
```



# Marshaling ICMPv6 messages

- **Simplicity wins**; allocating is okay until your performance needs are not met

```
im := icmp.Message{  
    Type: m.Type(),  
    Body: &icmp.DefaultMessageBody{  
        Data: mb,  
    },  
}
```

```
return im.Marshal(nil)
```



# ndp.NeighborSolicitation marshaling

- **Validate** before you **allocate**

```
func (ns *NeighborSolicitation) marshal() ([]byte, error) {  
    // Validation  
  
    // Allocation  
}
```



# Validate before you allocate

- Don't bother allocating memory until you've checked your inputs

```
// Only accept IPv6 target.  
if err := checkIPv6(ns.TargetAddress); err != nil {  
    return nil, err  
}
```



# Allocate once, if possible

- Allocating once is ideal for speed, but keep it simple

```
// Allocate enough space for base message.
```

```
b := make([]byte, nsLen)
```

```
copy(b[4:], ns.TargetAddress)
```

```
// Append any option bytes.
```


```
ob, err := marshalOptions(ns.Options)
```

```
if err != nil {
```

```
    return nil, err
```

```
}
```

```
b = append(b, ob...)
```



# When allocating memory...

- **Simplicity wins**; allocating is okay!
- **Write comprehensive unit tests** to lock in your behavior
- **Measure** for bottlenecks using Go benchmarks and pprof
- **Optimize** only after finding evidence of performance issues



# ndp.Message API




## ndp.Message types

<b>NDP RFC messages</b>	<b><a href="https://github.com/mdlayher/ndp">github.com/mdlayher/ndp</a></b>
Neighbor Advertisement	<code>ndp.NeighborAdvertisement</code>
Neighbor Solicitation	<code>ndp.NeighborSolicitation</code>
Router Advertisement	<code>ndp.RouterAdvertisement</code>
Router Solicitation	<code>ndp.RouterSolicitation</code>
Redirect	n/a, not necessary yet

## ndp.Message usage

```
m := &ndp.NeighborSolicitation{
    TargetAddress: target,
    Options: []ndp.Option{&ndp.LinkLayerAddress{
        Direction: ndp.Source,
        Addr:      addr,
    }},
}

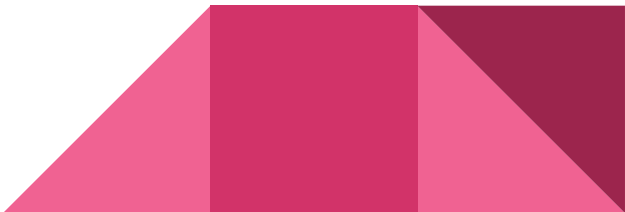
b, err := ndp.MarshalMessage(m)
if err != nil {
    return fmt.Errorf("failed to marshal: %v", err)
}
```



# ndp.Message usage

```
m, err := ndp.ParseMessage(b[:n])
if err != nil {
    return fmt.Errorf("failed to parse: %v", err)
}
```

```
switch m := m.(type) {
case *ndp.NeighborAdvertisement:
    printNA(m)
case *ndp.NeighborSolicitation:
    printNS(m)
default:
    log.Printf("%#v", m)
}
```





From bytes to options

# NDP option basics

- Options are encoded in type, length, value format
  - Fixed length: type
  - Fixed length: length
  - Variable length: value/data




# TLV options

Type (8 bits)	Length (8 bits)	Data (N bytes)
...		



# ndp.Option interface

```
// An Option is a Neighbor Discovery Protocol option.  
type Option interface {  
    // Code specifies the NDP option code for an Option.  
    Code() uint8  
  
    // Called when dealing with a Message's Options.  
    marshal() ([]byte, error)  
    unmarshal(b []byte) error  
}
```



# Parsing options

- An NDP message will always precede options
- NDP options on their own are not useful without an NDP message
- Exporting marshal/unmarshal methods bloats the API and GoDoc
  - **Solution:** use unexported functions with `ndp.ParseMessage` and `ndp.MarshalMessage`





# marshalOptions function

```
// marshalOptions marshals Options into a single byte slice.  
func marshalOptions(options []Option) ([]byte, error) {  
    // For each option...  
  
        // Marshal the option  
  
        // Append it to the output  
}
```



# parseOptions function

// parseOptions parses a slice of Options from a byte slice.

```
func parseOptions(b []byte) ([]Option, error) {
```

```
    // Iterate until no bytes remain...
```

```
        // Bounds check!!!
```

```
        // Read 2 bytes: type/length
```

```
        // Determine if option is known
```

```
        // Append to output slice
```

```
}
```



# ndp.Option types

- Common option types:
  - Source/target link-layer address
  - MTU
  - Prefix information
  - Recursive DNS server
  - ... and more! API could bloat quickly!



# Tips for implementing options

- Consider only implementing the most common options in your package
  - Prevent API bloat, support 90% of use cases
- **Tip:** add a “raw option” type or similar to enable further extension



# ndp.RawOption type

// A RawOption is an Option in its raw and unprocessed format.

// Unknown Options can be represented using a RawOption.

```
type RawOption struct {  
    Type    uint8  
    Length  uint8  
    Value   []byte  
}
```

// Code implements Option.

```
func (r *RawOption) Code() byte { return r.Type }
```




## ndp.Option types

NDP RFC options	<a href="https://github.com/mdlayher/ndp">github.com/mdlayher/ndp</a>
Source/Target link-layer address	<code>ndp.LinkLayerAddress</code>
MTU	<code>ndp.MTU</code>
Prefix information	<code>ndp.PrefixInformation</code>
Recursive DNS server (RDNSS)	<code>ndp.RecursiveDNSServer</code>
???	<code>ndp.RawOption</code>

# ndp.Option usage

```
var ra ndp.RouterAdvertisement
ra.Options = []ndp.Option{
    &ndp.LinkLayerAddress{
        Direction: ndp.Source,
        Addr:      addr,
    },
    ndp.NewMTU(1500),
    &ndp.PrefixInformation{
        PrefixLength: 32,
        Prefix:       net.ParseIP("2001:db8::"),
        SLAAC:        true,
    },
}
```



# Fuzzing byte parsers



# panic: runtime error: slice bounds out of range

```
goroutine 127 [running]:
testing.tRunner.func1(0xc4201453b0)
    /usr/local/go/src/testing/testing.go:742 +0x29d
panic(0x5dd240, 0x7378d0)
    /usr/local/go/src/runtime/panic.go:502 +0x229
github.com/mdlayher/ndp.(*NeighborAdvertisement).unmarshal(0xc42013d280, 0xc4200dd924,
0x10, 0x1c, 0xc4200dd920, 0x10)
    /home/matt/src/github.com/mdlayher/ndp/message.go:149 +0x2b1
github.com/mdlayher/ndp.ParseMessage(0xc4200dd920, 0x14, 0x20, 0x4, 0x14, 0xc4200dd920,
0x4)
    /home/matt/src/github.com/mdlayher/ndp/message.go:85 +0x168
github.com/mdlayher/ndp_test.TestParseMessageError.func1.1(0xc4201453b0)
    /home/matt/src/github.com/mdlayher/ndp/message_test.go:176 +0xd3
testing.tRunner(0xc4201453b0, 0xc420141440)
    /usr/local/go/src/testing/testing.go:777 +0xd0
created by testing.(*T).Run
    /usr/local/go/src/testing/testing.go:824 +0x2e0
```

# Enter go-fuzz

- If you're parsing raw bytes, there's a high potential for unexpected behavior:
  - Bad input causing application problems
  - Possibility of a panic taking down your program!
- [github.com/dvyukov/go-fuzz](https://github.com/dvyukov/go-fuzz)
  - Throws arbitrary bytes at your parser and finds crashers!
  - Mark inputs as “interesting” or not to guide fuzzer



# go-fuzz setup

```
//+build gofuzz
```

```
package ndp
```

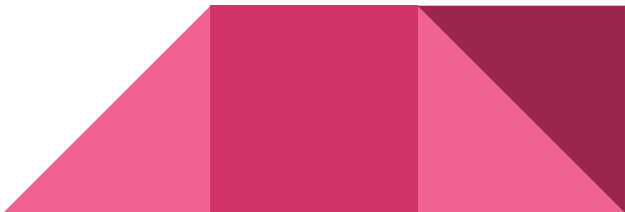
```
// Fuzz is an entry point for go-fuzz.
```

```
func Fuzz(data []byte) int {  
    return fuzz(data)  
}
```



# go-fuzz setup

```
func fuzz(data []byte) int {  
    m, err := ParseMessage(data)  
    if err != nil {  
        return 0 // Invalid, not interesting!  
    }  
    b2, err := MarshalMessage(m)  
    if err != nil {  
        panic(err)  
    }  
    if _, err := ParseMessage(b2); err != nil {  
        panic(err)  
    }  
    return 1 // Valid, interesting!  
}
```



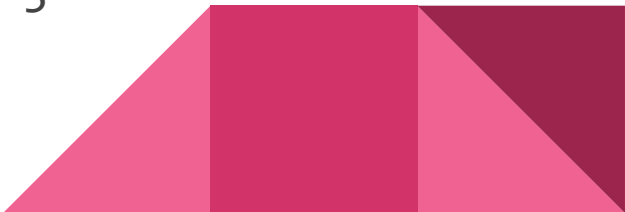
# go-fuzz usage

- Prepare the fuzzer by building an instrumented test program

```
$ CGO_ENABLED=0 go-fuzz-build github.com/mdlayher/ndp
```

- Run go-fuzz with multiple CPUs and output results to ./fuzz/

```
$ go-fuzz -bin ./ndp-fuzz.zip -procs 16 -workdir ./fuzz/  
... workers: 16, corpus: 78 (0s ago), crashers: 5  
^C
```



## go-fuzz usage

- Inspect the resulting crasher inputs

```
$ cat fuzz/crashers/4c24217a9963fae05ea48d657c342549a731989.quoted  
"\x86000000000000000000\x05"
```

- Write a test, fix the bug, and repeat!

```
fuzz([]byte("\x860000000000000000\x05"))
```

# go-fuzz conclusions (use it!)

- Use go-fuzz on ALL byte parsers: [github.com/dvyukov/go-fuzz](https://github.com/dvyukov/go-fuzz)
  - Find parsing problems now, not during a 3am outage
  - A multi-worker mode is available for use with clusters of machines
  - Submit “trophies” to the go-fuzz README





ndp.Conn struct



# Building connection types

- “Conn” types represent network connections
  - “Dial” and/or “Listen” constructors
  - “Close” to free resources
  - “Read” and “Write” to pass messages



# net vs x/net

- NDP is transported over IPv6 + ICMPv6
- Standard library net doesn't quite provide all the functionality we need
  - [golang.org/x/net](https://golang.org/x/net) is designed for advanced use-cases!



# ICMPv6 networking in Go

- [golang.org/x/net/icmp](https://golang.org/x/net/icmp)
- [golang.org/x/net/ipv6](https://golang.org/x/net/ipv6)
- Huge shout-out to [Mikio Hara](#) for his work on these packages and countless other low-level networking packages for Go



# ICMPv6 listener

- Privileged operation; usually needs root

```
// Open raw ICMPv6 listener on eth0's link-local address.  
addr := "fe80::7d64:35ff:fee7:cbc4%eth0"  
ic, err := icmp.ListenPacket("ip6:ipv6-icmp", addr)  
if err != nil {  
    return err  
}
```



# Reading ICMPv6 messages

- Similar to standard APIs, but also returns IPv6 control messages

```
b := make([]byte, 1024)
n, cm, src, err := c.pc.ReadFrom(b)
if err != nil {
    return nil, nil, nil, err
}

return b[:n], cm, src.IP, nil
```



# Writing ICMPv6 messages

- Similar to standard APIs, but you can specify IPv6 control messages

```
// Write bytes to the specified target.  
_, err := c.pc.WriteTo(b, cm, &net.IPAddr{  
    IP:    ip,  
    Zone: c.ifc.Name,  
})  
return err
```






ndp.Conn usage

# Creating an `ndp.Conn`

- Select an interface, dial ICMPv6, specify an address to listen on

```
ifi, err := net.InterfaceByName("eth0")
if err != nil {
    log.Fatalf("failed to get interface: %v", err)
}
```

```
// Dial IPv6 + ICMPv6 connection.
c, ip, err := ndp.Dial(ifi, ndp.LinkLocal)
if err != nil {
    log.Fatalf("failed to dial NDP: %v", err)
}
```





# Reading `ndp.Messages`

- Keep reading and printing messages until an error occurs

```
for {  
    msg, _, from, err := c.ReadFrom()  
    if err != nil {  
        return nil, err  
    }  
  
    printMessage(msg, from)  
}
```

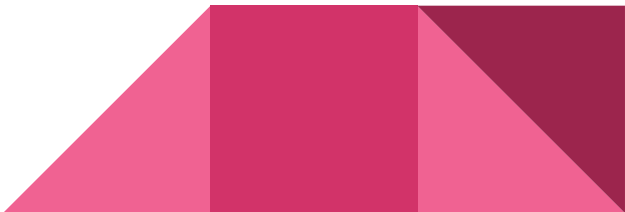


# Writing ndp.Messages

- Send a router solicitation to trigger router advertisements on the network

```
m := &ndp.RouterSolicitation{
    Options: []ndp.Option{&ndp.LinkLayerAddress{
        Direction: ndp.Source, Addr: addr,
    }},
}
```

```
dst := net.IPv6linklocalallrouters
if err := c.WriteTo(m, nil, dst); err != nil {
    return nil, err
}
```





Build a tool to test your package

# Build a tool to test your package

- Add a cmd/ directory with a testing utility during development
- Consider building it out to become a useful tool
  - [cmd/ndp](#): tool for generating and capturing NDP traffic



# Introducing the ndp tool

- `$ ndp [listen]`
  - Listen for any NDP messages that pass through the interface
- `$ ndp rs`
  - Send a router solicitation; wait for a router advertisement
- `$ ndp -t fd00::1 ns`
  - Send a neighbor solicitation; wait for a neighbor advertisement





tcpdump?

```
$ sudo tcpdump -i enp4s0 'icmp6 && (ip6[40] == 133 or ip6[40] == 134)'  
tcpdump: verbose output suppressed, use -v or -vv for full protocol  
decode  
listening on enp4s0, link-type EN10MB (Ethernet), capture size 262144  
bytes  
16:02:42.774725 IP6 nerr-2 > ip6-allrouters: ICMP6, router  
solicitation, length 16  
16:02:42.777116 IP6 _gateway > ip6-allnodes: ICMP6, router  
advertisement, length 88
```

ndp!



```
$ sudo ./bin/ndp rs
ndp> interface: enp4s0, link-layer address: 74:d4:35:e7:cb:c4, IPv6
address: fe80::e563:9887:3aca:e01e
ndp rs> router solicitation:
    - source link-layer address: 74:d4:35:e7:cb:c4

ndp rs> router advertisement from: fe80::618:d6ff:fea1:ceb7:
    - hop limit:          64
    - router lifetime:    30m0s
    - options:
        - prefix information: 2600:6c4a:787f:d200::/64, flags: [0A],
valid: 24h0m0s, preferred: 4h0m0s
        - prefix information: fd00::/64, flags: [0A], valid: 24h0m0s,
preferred: 4h0m0s
        - source link-layer address: 04:18:d6:a1:ce:b7
```

# Troubleshooting your ISP's equipment with Go

- IPv6 works for a few days...
- Ubiquiti EdgeRouter Lite can run Go programs!

```
desktop $ GOARCH=mips64 go build -o ndp_mips64
```

```
desktop $ scp ndp_mips64 router:~/ndp
```


```
router $ sudo ./ndp -i eth1 rs
```



```
$ sudo ./ndp -i eth1 rs
ndp> interface: eth1, link-layer address: 04:18:d6:a1:ce:b7, IPv6
address: fe80::618:d6ff:fea1:ceb7
ndp rs> router solicitation:
    - source link-layer address: 04:18:d6:a1:ce:b7
.....
.....^C
ndp rs> sent 95 router solicitation(s)
```

# Troubleshooting your ISP's equipment with Go

- No luck with tech support: “your WiFi router isn’t working”
- ... a modem swap during an upgrade made the problem disappear



# Conclusions

# Conclusions

- IPv6 is great, check out [ipv6-test.com](https://ipv6-test.com) to see if you're using it
- Network protocols are powerful building blocks
- Go is an excellent language for exploring low-level network protocols
- Build tools to solve real problems on your network!



# Resources

- [github.com/mdlayher/ndp](https://github.com/mdlayher/ndp)
- [Network Protocol Breakdown: NDP and Go](#)
- [RFC 4861](#)



# Thanks!

**Matt Layher**

[github.com/mdlayher](https://github.com/mdlayher)

[twitter.com/mdlayher](https://twitter.com/mdlayher)

Image credit: [worldipv6launch.org](https://worldipv6launch.org)

