

Using `unsafe.Pointer` to explore Linux system calls

Matt Layher

GoCon Canada, May 31, 2019

Matt Layher

- GitHub, Twitter: **@mdlayher**
- github.com/mdlayher/talks



fastly[®]

unsafe

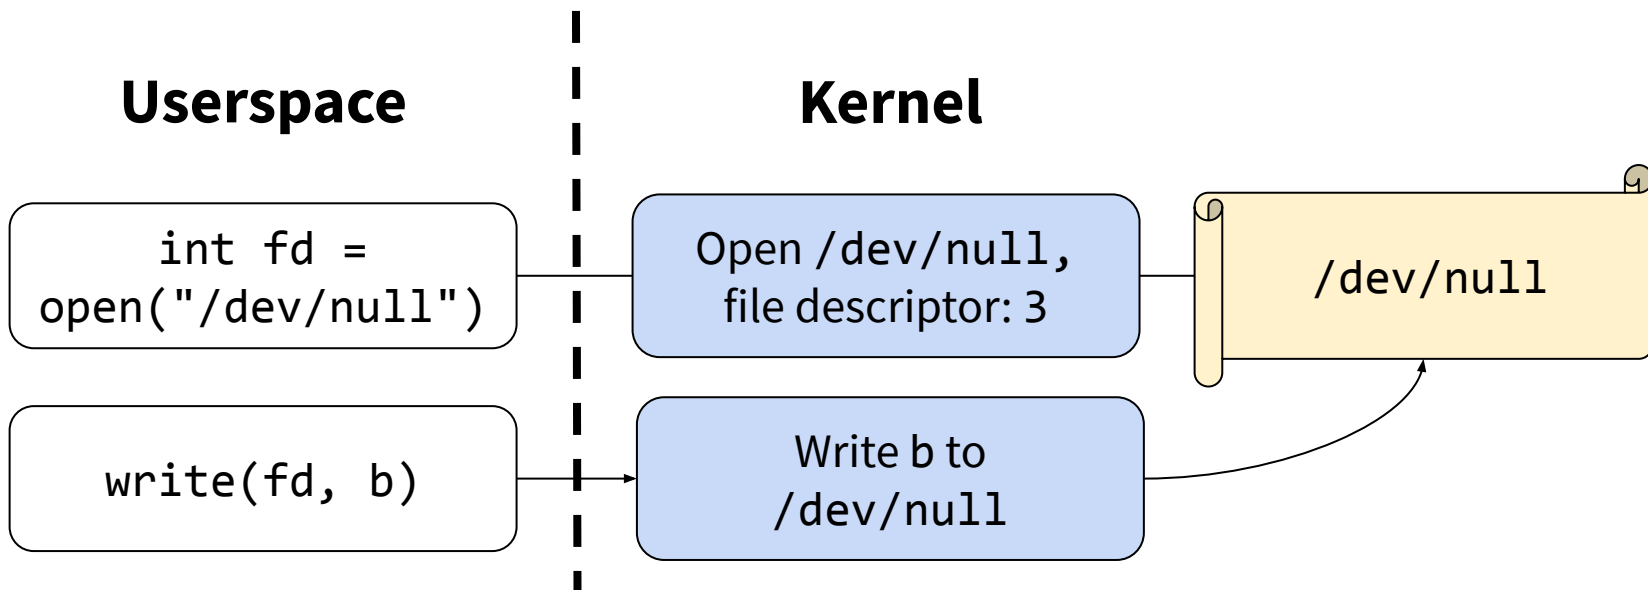
Package unsafe contains operations that step around the type safety of Go programs.

unsafe and Linux system calls

- Package unsafe and Linux **system calls** go hand-in-hand
- Sometimes an escape hatch from Go's type system is necessary
- **Read the rules before you write any unsafe code**
 - golang.org/pkg/unsafe/#Pointer
 - `go vet` can catch some mistakes, but don't rely on it

What is a system call?

- A “function call” into the Linux kernel, used to access to files, hardware, etc.



ioctl(2) on Linux

- Short for “I/O control”
- “a catch-all for operations that don't cleanly fit the UNIX stream I/O model”
 - UNIX’s “everything is a file” model: `read(2)` and `write(2)`
- Primarily used to pass data structures between the kernel and userspace

Retrieving data with `ioctl(2)` in C

- Suppose we want to retrieve the VM sockets context ID for our system

```
int fd = open("/dev/vsock", O_RDONLY);
if (fd < 0) {
    perror("failed to open file");
}

uint32_t cid;
if (ioctl(fd, IOCTL_VM_SOCKETS_GET_LOCAL_CID, &cid) < 0) {
    perror("failed to get local CID");
}
printf("CID: %d\n", cid);
```

**How can we
replicate this
program in Go?**

Prerequisites

- Discuss type safety, memory layout, and endianness
- Introduce `unsafe` and explain its use cases
- Establish some guidelines on how to make **safe** use of `unsafe`

What is type safety?

- **Type safety** provides safeguards for the programmer, preventing mistakes
- Go is **statically typed**: the compiler checks data types and enforces type safety

```
fmt.Println(1 + "abc")
```

```
// BAD: cannot convert "abc" (type untyped string) to type int
```

```
// BAD: invalid operation: 1 + "abc" (mismatched types int and string)
```

```
fmt.Println(strconv.Itoa(1) + "abc")
```

```
// OK: 1abc
```


Memory layout of Go types

- `intN/uintN` family are fixed size integers, 8 bits per byte

`int8, uint8`




`int16, uint16`



`int32, uint32`




`int64, uint64`



Memory layout of Go types

- `int/uint` size varies by CPU architecture; **avoid them** when using unsafe

`int, uint (386)`



`int, uint (amd64)`




Memory layout of Go types

- The same rules apply for **arrays**, but what about **slices**?

[1]byte 


[2]byte 

[4]byte 

[8]byte 

Memory layout of Go types

- Slices must be handled carefully with unsafe code; **remember the slice header!**




blog.golang.org/go-slices-usage-and-internals

What is endianness?


- The way a particular CPU lays out values in memory (also called **byte order**)

`uint16(1): big endian`



Network byte order

`uint16(1): little endian`



x86 CPUs (Intel, AMD)

What is endianness?

- You can store integers as either big or little endian with encoding/binary

```
v := uint16(1)
big := make([]byte, 2)
little := make([]byte, 2)
```

```
binary.BigEndian.PutUint16(big, v)
binary.LittleEndian.PutUint16(little, v)
```

```
fmt.Println(big, little)
// OK: [0 1] [1 0]
```

unsafe

```
$ go doc unsafe  
package unsafe // import "unsafe"
```

Package unsafe contains operations that step around the type safety of Go programs.

Packages that import unsafe may be non-portable and are not protected by the Go 1 compatibility guidelines.

```
func Alignof(x ArbitraryType) uintptr  
func Offsetof(x ArbitraryType) uintptr  
func Sizeof(x ArbitraryType) uintptr  
type ArbitraryType int  
type Pointer *ArbitraryType
```

unsafe.Sizeof()

unsafe.Sizeof() with integers (amd64)

- How much memory does a value actually occupy?

```
const (  
    size8      = unsafe.Sizeof(uint8(0))  
    size16     = unsafe.Sizeof(uint16(0))  
    size32     = unsafe.Sizeof(uint32(0))  
    size64     = unsafe.Sizeof(uint64(0))  
    sizeUint   = unsafe.Sizeof(uint(0))  
)
```

```
fmt.Println(size8, size16, size32, size64, sizeUint)  
// 1 2 4 8 8
```

unsafe.Sizeof() with a struct (amd64)

- Shouldn't this struct occupy 14 bytes?

```
var s struct {  
    One    uint64 // 8  
    Two    uint32 // 4  
    Three  uint16 // 2  
}
```

```
fmt.Printf("want: %d, got: %d", 8+4+2, unsafe.Sizeof(s))  
// want: 14, got: 16
```

unsafe.Sizeof() with a padded struct (amd64)

- Struct definitions are padded to the next **machine word size (64 bits)**

```
var s struct {  
    One    uint64 // 8  
    Two    uint32 // 4  
    Three  uint16 // 2  
    _      [2]byte // 2  
}
```

```
fmt.Printf("want: %d, got: %d", 8+4+2+2, unsafe.Sizeof(s))  
// want: 16, got: 16
```


unsafe.Pointer

Provided that T2 is no larger than T1 and that the two share an equivalent memory layout, `[unsafe.Pointer]` conversion allows reinterpreting data of one type as data of another type.

golang.org/pkg/unsafe/#Pointer

Converting uint16 to [2]byte

- Go's type system won't allow this conversion

```
a := uint16(1)
```

```
// BAD: cannot convert a (type uint16) to type [2]byte
```

```
b := [2]byte(a)
```

Converting uint16 to [2]byte

- `unsafe.Pointer` conversions defeat Go's type system

```
// Always check sizes before performing these conversions!
a := uint16(1)
if unsafe.Sizeof(a) != 2 {
    panic("a is not of the expected size")
}

b := *(*[2]byte)(unsafe.Pointer(&a))
fmt.Println(b)
// [1 0]
```

What does this actually do?

```
*( *[2]byte )(unsafe.Pointer(&a))
```

Take the address of a, producing a *uint16

```
*( *[2]byte )(unsafe.Pointer(&a))
```

Convert *uint16 to unsafe.Pointer

```
*(*[2]byte)(unsafe.Pointer(&a))
```

Convert `unsafe.Pointer` to `*[2]byte`

```
*(*[2]byte)(unsafe.Pointer(&a))
```


Dereference pointer, producing [2]byte

```
 (*[2]byte)(unsafe.Pointer(&a))
```

Breaking down the conversion

- You could write the same operation with intermediate variables if you wanted

```
a := uint16(1)
```

```
uint16Ptr := &a
```

```
unsafePtr := unsafe.Pointer(uint16Ptr)
```

```
arrayPtr := (*[2]byte)(unsafePtr)
```

```
b := *arrayPtr
```

```
fmt.Println(b)
```

```
// [1 0]
```

A note on slices versus arrays

- **Arrays** are generally used in unsafe conversions, **not slices**
 - Slice an array after conversion, or take address of first element of slice

```
a := uint16(1)
```

```
b := (*( *[2]byte)(unsafe.Pointer(&a)) )[:]
```

```
fmt.Println(b)
```

```
// [1 0]
```

```
fmt.Println(*( *uint16)(unsafe.Pointer(&b[0])))
```

```
// 1
```

The danger zone

The danger zone

- The `unsafe.Pointer` documentation covers 6 patterns which are valid
 - **Read and understand these patterns before you make use of `unsafe`!**



Credit: Ashley McNamara, github.com/ashleymcnamara/gophers

The danger zone: reading arbitrary memory

- You **must** be judicious in your use of unsafe: **always check type sizes**

```
// All bets are off on what is actually stored in b.
```

```
a := uint16(1)
```

```
b := *(*[4]byte)(unsafe.Pointer(&a))
```

```
fmt.Println(b)
```

```
// [1 0 160 93]  <- 1st run in play.golang.org
```

```
// [1 0 168 118] <- 2nd
```

```
// [1 0 19 119]  <- 3rd
```

The danger zone: pointer arithmetic

- Remember this from C? Say hello to uintptr!
 - Take the address of the first element in the array, add ($i * 4$) each iteration

```
b := []uint32{1, 2, 3, 4}
for i := 0; i < len(b); i++ {
    // Pretty much the worst possible way to print a slice in Go.
    fmt.Printf("%d ", *(*uint32)(unsafe.Pointer(uintptr(
        unsafe.Pointer(&b[0])) + uintptr(i)*unsafe.Sizeof(b[0])),
    ))
}
// 1 2 3 4
```

The danger zone: writing arbitrary memory

- Even worse, you can **write to arbitrary memory**

```
var v uint32
```

```
// Overwrite whatever data lives at this address.
```

```
((*[4]byte)(  
    unsafe.Pointer(uintptr(unsafe.Pointer(&v)) - 0xffffffff),  
)) = [4]byte{0xff, 0xff, 0xff, 0xff}
```


unexpected fault address 0xbf0004c777

fatal error: fault

[signal SIGSEGV: segmentation violation code=0x1 addr=0xbf0004c777 pc=0x487268]

goroutine 1 [running]:

runtime.throw(0x4b8e08, 0x5)

 /usr/local/go/src/runtime/panic.go:617 +0x72 fp=0xc00004c700

sp=0xc00004c6d0 pc=0x427fc2

runtime.sigpanic()

 /usr/local/go/src/runtime/signal_unix.go:397 +0x401 fp=0xc00004c730

sp=0xc00004c700 pc=0x43a8f1

main.main()

 /home/matt/src/github.com/mdlayher/tmp/main.go:12 +0x38 fp=0xc00004c798

sp=0xc00004c730 pc=0x487268

runtime.main()

 /usr/local/go/src/runtime/proc.go:200 +0x20c fp=0xc00004c7e0

sp=0xc00004c798 pc=0x42992c

runtime.goexit()

 /usr/local/go/src/runtime/asm_amd64.s:1337 +0x1 fp=0xc00004c7e8

sp=0xc00004c7e0 pc=0x4511b1

exit status 2

**When is unsafe
actually
appropriate to use?**

~~syscall~~

golang.org/x/sys/unix

unsafe with Linux system calls

- Working with system calls often involves unsafe operations
- `syscall` is deprecated in favor of `golang.org/x/sys`
 - golang.org/s/go1.4-syscall
- `ioctl(2)`, `getsockopt(2)`, `setsockopt(2)`
 - These system calls are too flexible to expose a general-purpose API
 - `unix.IoctlGetInt()`, `unix.SetsockoptLinger()`, etc.

ioctl(2) in x/sys/unix

- We can pass a pointer (integer memory address) or even just a regular integer

```
func ioctl(fd int, req uint, arg uintptr) (err error) {  
    _, _, e1 := Syscall(SYS_IOCTL,  
        uintptr(fd), uintptr(req), uintptr(arg),  
    )  
    if e1 != 0 {  
        err = errnoErr(e1)  
    }  
    return  
}
```

Retrieving data with `ioctl(2)` in C

```
int fd = open("/dev/vsock", O_RDONLY);
if (fd < 0) {
    perror("failed to open file");
}

uint32_t cid;
if (ioctl(fd, IOCTL_VM_SOCKETS_GET_LOCAL_CID, &cid) < 0) {
    perror("failed to get local CID");
}
printf("CID: %d\n", cid);
```

If [an `unsafe.Pointer`] argument must be converted to `uintptr` for use as an argument, that conversion must appear in the call expression itself.

golang.org/pkg/unsafe/#Pointer

Retrieving data with ioctl(2) in Go

```
f, err := os.Open("/dev/vsock")
if err != nil {
    log.Fatalf("failed to open file: %v", err)
}

var cid uint32
if err := ioctl(int(f.Fd()), unix.IOCTL_VM_SOCKETS_GET_LOCAL_CID,
    uintptr(unsafe.Pointer(&cid)),
); err != nil {
    log.Fatalf("failed to get local CID: %v", err)
}
fmt.Printf("CID: %d\n", cid)
```


Native endianness

- Integers passed across the userspace/kernel boundary use **native endianness**

```
func nativeEndian() binary.ByteOrder {  
    a := uint16(1)  
    switch *(*[2]byte)(unsafe.Pointer(&a)) {  
    case [2]byte{0, 1}:  
        return binary.BigEndian  
    case [2]byte{1, 0}:  
        return binary.LittleEndian  
    default:  
        panic("unknown endianness")  
    }  
}
```

Linux taskstats interface

- A large type (**328 bytes**); parsing bytes into fields manually could be error prone

```
type Taskstats struct {  
    Version          uint16    Ac_ppid          uint32    Ac_utimescaled    uint64  
    Ac_exitcode       uint32    Ac_btime         uint32    Ac_stimescaled    uint64  
    Ac_flag           uint8     Ac_etime         uint64    Cpu_scaled_run_real_total uint64  
    Ac_nice           uint8     Ac_uptime        uint64    Freepages_count   uint64  
    Cpu_count         uint64    Ac_stime         uint64    Freepages_delay_total uint64  
    Cpu_delay_total   uint64    Ac_minflt        uint64    Thrashing_count   uint64  
    Blkio_count        uint64    Ac_majflt        uint64    Thrashing_delay_total uint64  
    Blkio_delay_total uint64    Coremem          uint64  
    Swapin_count      uint64    Virtmem          uint64  
    Swapin_delay_total uint64    Hiwater_rss      uint64  
    Cpu_run_real_total uint64    Hiwater_vm       uint64  
    Cpu_run_virtual_total uint64    Read_char        uint64  
    Ac_comm           [32]uint8 Write_char       uint64  
    Ac_sched          uint8     Read_syscalls    uint64  
    Ac_pad            [3]uint8  Write_syscalls   uint64  
    _                 [4]byte   Read_bytes       uint64  
    Ac_uid            uint32    Write_bytes      uint64  
    Ac_gid            uint32    Cancelled_write_bytes uint64  
    Ac_pid            uint32    Nvcsw            uint64  
                    Nivcsw            uint64  
}
```

Linux taskstats interface

- A single `unsafe.Pointer` conversion to `unix.Taskstats` is all we need

```
b := []byte{0x01, /* ... */}
const sizeofTaskstats = int(unsafe.Sizeof(unix.Taskstats{}))

// Always sanity check the structure size before conversion!
if sizeofTaskstats != len(b) {
    return nil, errors.New("unexpected taskstats structure size")
}

stats := (*unix.Taskstats)(unsafe.Pointer(&b[0]))
```

Other uses for unsafe

- Potential performance improvements in specific situations
- Cgo: passing data between C and Go
- Accessing unexported identifiers with `//go:linkname`

Summary

With great power comes great responsibility

- When you import unsafe, **you're expected to know how to use it safely**
- Do not fear unsafe, it's a vital part of what makes Go work
- When in doubt, seek guidance and ask questions!
 - Gophers Slack: invite.slack.golangbridge.org, #darkarts

Resources and thanks

- Blog: `unsafe.Pointer` and system calls
 - mdlayher.com/blog/unsafe-pointer-and-system-calls
- Source for packages referenced in this talk
 - github.com/mdlayher/taskstats, github.com/mdlayher/vsock
- Thanks to @acln, @jadr2ddude, @kale, @pwaller, @seeb, and @zeebo from #darkarts on Gophers Slack for their review

Thanks!

Matt Layher

mdlayher.com

github.com/mdlayher

twitter.com/mdlayher

Go gopher by Renee French, licensed under Creative Commons 3.0 Attributions license.

